

From Predictive Coverage to Robust PDE Control: Calibrated Dynamics Ambiguity Sets for Neural Operators

Yaowen Kang

Department of Mathematics, Imperial College London

Advisor: Nicolas Boullé

Abstract

Neural PDE world models can be accurate in an average field norm yet err along the directions that determine a control action. Predictive uncertainty becomes useful for control only when its geometry can be queried by a planner. We construct function-valued dynamics ambiguity sets from a residual Fourier neural operator (FNO) and a replica trained on Gaussian-perturbed labels. Split conformal calibration turns their smoothed disagreement into an isotropic L^2 tube, a normalized ellipsoid, or a simultaneous coordinate box. Robust model predictive control queries these sets through exact adjoint support functions or projected-gradient adversarial rollouts. The analysis separates fixed-grid marginal coverage from deterministic rollout and value bounds that require a uniform one-step error premise. Under compound shift in controlled Burgers dynamics, the audit ellipsoid and simultaneous box both attain 0.906 field coverage, yet their adjoint controllers produce empirical $p90$ costs of 1.0201 and 0.9586. The latter is 10.84% below nominal MPC across 24 matched actuator-gain cases. On a two-dimensional Navier–Stokes benchmark, simultaneous coverage is 0.883 at a nominal 0.90 level, but the bands remain wide because the perturbation scale weakly localizes error. Calibration determines predictive validity; ambiguity geometry and the planner query determine whether that validity becomes useful for control.

Contents

1	Introduction	3
1.1	Related work	3
1.2	Contributions	4
2	Preliminaries and problem setting	4
2.1	Learned dynamics on discretized PDE fields	4
2.2	Finite-horizon control objective	5
2.3	Split conformal prediction on a discretized field	5
3	Data preparation	5
3.1	Controlled Burgers data generation	5
3.2	Public two-dimensional Navier–Stokes data	6
3.3	Discretized benchmark dynamics and learned surrogates	7
3.4	Data roles and split protocol	7
4	Methodology	8
4.1	Control-oriented dynamics ambiguity sets	8
4.2	Residual FNO world model	8
4.3	Perturbation-based uncertainty scale	8
4.4	Perturbation-based nonconformity scores and calibration	8
4.5	Practical implementation	9
4.6	From one-step calibrated sets to multi-step robust control	9
4.6.1	Exact support and the adjoint robust counterpart	10
4.6.2	Nonlinear adversarial rollout	10
4.6.3	Outer optimization and receding-horizon deployment	11
4.7	Theoretical guarantees and scope	11
4.7.1	Finite-sample coverage on a fixed discretization	11
4.7.2	First-order robustification and curvature	11
4.7.3	Multi-step propagation under uniform one-step error	11
4.7.4	Discounted value and policy transfer	12
4.7.5	Guarantee separation	13
5	Experimental evaluation	13
5.1	Baselines, metrics, and implementation	13
5.2	Ablation protocol and within-group controls	14
5.3	World-model fit and uncertainty localization	14
5.4	Ablation A: audit calibration restores coverage under shift	14
5.5	Ablation B: matched-coverage sets yield different tail costs	15
5.6	Navier–Stokes stress test: validity survives, localization weakens	16
5.7	Ablation C: tighter value bounds require matched coverage	18
5.8	Interpretation and limitations	19
6	Conclusion	20
	Data availability	20
	Code availability	20
	References	20

1 Introduction

Many PDE control problems require repeated solution of a nonlinear evolution equation while an optimizer searches over candidate actions. Neural operators amortize this computation by learning solution maps between function spaces [6, 17]. Fourier neural operators (FNOs) and DeepONets provide two widely used realizations and have been evaluated on Burgers, Darcy, and Navier–Stokes systems [20, 22]. This speed makes a learned operator a natural world model for model predictive control (MPC).

Mean prediction accuracy alone is not a decision guarantee. Neural-operator errors can change under input-distribution shift, even when in-domain tests are accurate [29]. Probabilistic, Bayesian, ensemble, and linearized operator models estimate structured predictive uncertainty [7, 25, 26]. These uncertainty estimates are useful diagnostics, but their dispersion does not by itself provide a distribution-free finite-sample coverage statement.

Conformal prediction (CP) provides a complementary calibration layer. Split CP wraps a fixed nonconformity score and, under exchangeability, yields marginal finite-sample coverage for a new response [1, 19, 36]. Recent work has adapted this principle to operator outputs and simultaneous fields [21, 23, 30, 37]. Coverage, however, only specifies a containment event. It does not determine how the geometry of the prediction set changes a control action.

This distinction is central in sequential decision making. Value-aware model learning shows that prediction errors should be judged through their effect on downstream value [12]. Lipschitz analyses show how one-step model discrepancies can propagate through multi-step predictions and value estimates [2]. Robust optimization and robust dynamic programming instead optimize against a declared uncertainty set [5, 15, 32]. The missing interface is therefore not another point predictor. It is a dynamics ambiguity set whose coverage is calibrated and whose geometry can be queried by the control objective.

We construct this interface for discretized PDE fields. A pair of FNO world models provides a perturbation-based local scale. A disjoint target-regime audit then calibrates several field-valued ambiguity geometries. Robust MPC queries each geometry through an adjoint support function or a nonlinear adversarial rollout. The construction separates predictive validity, set efficiency, and closed-loop utility, since none of these properties implies the others.

We develop that interface through the chain

$$(\hat{G}, \tilde{G}) \longrightarrow \sigma(x, a) \xrightarrow{\text{split CP}} \mathcal{U}_g(x, a) \xrightarrow{\text{support or adversary}} \text{robust MPC}.$$

Sensitivity to label perturbations supplies the local scale σ ; a disjoint deployment audit calibrates the geometry g ; and the finite-horizon objective queries the resulting field set. The FNO is a replaceable world model, not the architectural contribution. The current planner uses a stagewise trajectory model whose scope remains separate from one-step conformal coverage.

1.1 Related work

Operator learning. Neural operators formalize parameter-shared maps between function spaces and support evaluation across discretizations [17]. FNOs parameterize global integral kernels in Fourier space [20], while DeepONet uses a branch–trunk representation grounded in operator approximation [22]. We use an FNO only as the differentiable one-step surrogate; the proposed uncertainty-to-decision interface is architecture agnostic.

Uncertainty quantification. Operator-specific UQ methods include ensembles for out-of-domain PDE inputs [29], probabilistic neural operators [7], approximate Bayesian neural operators [25], and linearized function-valued Gaussian processes [26]. These approaches model predictive dispersion or a distribution over output functions. Our method instead uses perturbation sensitivity only as a local scale and delegates finite-sample validity to a separate calibration split.

Conformal prediction. Classical CP obtains finite-sample marginal validity through exchangeable ranks [1, 19, 36]. For operator outputs, risk-controlling quantile neural operators calibrate functional risk [23], Conformalized-DeepONet constructs distribution-free bands [30], and split CP has been extended through a discretize–then–lift argument [21]. Most closely, Wang et al. use the disagreement of clean-label and perturbed-label FNOs in a max-type score for simultaneous Navier–Stokes fields [37]. We retain that scale construction but study how alternative calibrated geometries enter a controller.

Robust decisions and control. Conformal sets can serve as uncertainty sets in robust optimization [16]. Calibrated model-based reinforcement learning studies how uncertainty calibration affects planning [27], while distributionally robust model-based reinforcement learning optimizes against dynamics ambiguity [34]. Chee et

al. combine weighted CP with dynamic constraint tightening and reference generation [8]; adaptive CP has also been used for motion planning [11]. SODA-MPC uses conformalized ensembles with runtime adaptation [9], and conformalized system-level synthesis constructs OOD reachable tubes [33]. SafeDiffCon incorporates a conformal uncertainty quantile into diffusion-based PDE control [14]. These methods target constraints, reachable tubes, or generative control sequences. Our setting asks how a calibrated PDE-field set should be shaped and queried by a finite-horizon objective. We do not inherit the feasibility or safety guarantees of the cited controllers.

1.2 Contributions

The paper makes four scoped contributions.

- (i) A disjoint audit protocol converts perturbation sensitivity into function-space dynamics sets. Clean and perturbed FNOs use the full proper-training input set, while no calibration label trains either predictor.
- (ii) An objective-sensitive query converts those sets into decisions. We derive exact support functions for normalized L^2 ellipsoids and coordinate boxes, embed them in an adjoint robust counterpart, and compare this approximation with nonlinear adversarial rollouts.
- (iii) A guarantee ledger separates finite-sample marginal coverage, deterministic multi-step and value propagation under uniform error, and a conditional curvature remainder. None is presented as a substitute for another.
- (iv) Matched ablations separate calibration source, ambiguity geometry, and the robust inner query. Controlled Burgers supplies conditional closed-loop evidence, while Navier–Stokes exposes a high-dimensional localization failure. Negative results are retained as evidence about the interface, rather than removed as unsuccessful variants.




Figure 1: Closed-loop uncertainty-to-control architecture. Offline, the data roles in Section 3.4 feed the twin FNO, perturbation scale, and conformal calibration in Sections 4.2–4.5. Online, the calibrated field set is queried by the robust planner in Section 4.6. The selected action is applied to the controlled Burgers plant in Section 3.1, and the next observed field closes the receding-horizon loop. Dashed arrows denote offline transfer, not a claim that one-step marginal conformal coverage certifies the counterfactual closed loop.

2 Preliminaries and problem setting

2.1 Learned dynamics on discretized PDE fields

Neural operators approximate maps between function spaces before numerical evaluation on a grid [6, 17]. A trained implementation nevertheless receives finite arrays. We therefore distinguish the continuous solution operator from the discretized dynamics map used by the world model. All conformal statements in this paper concern the full output field on one declared grid. They do not imply continuum coverage or resolution transfer.

For a discretized field $v, w \in \mathbb{R}^n$, we use the quadrature-normalized inner product and norm

$$\langle v, w \rangle_n = \frac{1}{n} \sum_{j=1}^n v_j w_j, \quad \|v\|_{2,n}^2 = \langle v, v \rangle_n. \quad (1)$$

Nonuniform quadrature weights can replace $1/n$. This normalization keeps field errors comparable across the two benchmark dimensions while preserving the geometry of each declared score.

2.2 Finite-horizon control objective

For an action-conditioned learned transition $\hat{G}_\theta$, the finite-horizon objective is

$$J_H(x_0, a_{0:H-1}) = \sum_{t=1}^H \ell(x_t) + \sum_{t=0}^{H-1} r(a_t), \quad x_{t+1} = \hat{G}_\theta(x_t, a_t), \quad (2)$$

where lower cost is preferred. MPC optimizes the open-loop sequence, applies its first action, observes the next field, and replans. This follows the standard receding-horizon construction [28]. The optimizer queries states and actions that need not follow the one-step training distribution, which makes the treatment of dynamics uncertainty consequential for control.

2.3 Split conformal prediction on a discretized field

We use the standard split-conformal construction [1, 19]. Let $Z_i = (X_i, A_i, Y_i)$ and let $S(Z_i)$ be a scalar nonconformity score fixed before calibration. With calibration scores $S_{(1)} \leq \dots \leq S_{(m)}$, target miscoverage $\alpha \in (0, 1)$, and

$$k = \lceil (m+1)(1-\alpha) \rceil, \quad q_{1-\alpha} = S_{(k)}, \quad (3)$$

the standard convention sets $q_{1-\alpha} = +\infty$ when $k > m$. For a new input (x, a) , the prediction set is

$$C(x, a) = \{y : S(x, a, y) \leq q_{1-\alpha}\}. \quad (4)$$

Under exchangeability, the rank of the new score among the $m+1$ scores yields marginal coverage of at least $1-\alpha$ [19, 36]. This guarantee is not conditional coverage at every state-action pair. It also does not provide simultaneous coverage over a counterfactual MPC rollout [1].

3 Data preparation

The two benchmarks serve different purposes. Controlled Burgers supplies action-conditioned transitions and closed-loop trajectories generated by our solver. The public Navier–Stokes archive supplies a higher-dimensional test of field uncertainty calibration without an action channel.

3.1 Controlled Burgers data generation

The control benchmark is the one-dimensional viscous Burgers equation on $\Omega_B = (0, 1)$,

$$\partial_t u + u \partial_x u = \nu \partial_{xx} u + g_{\text{act}} a_t b(x), \quad u(t, 0) = b_L, \quad u(t, 1) = b_R, \quad (5)$$

where $\nu > 0$ is the viscosity and $a_t \in \mathcal{A} = [-2, 2]$ is held constant on $[t, t + \Delta T)$ with $\Delta T = 0.02$. The scalar g_{act} is the physical actuator gain. The localized actuator shape is parameterized by a center x_c and width s_b ,

$$b(x; x_c, s_b) = \frac{\exp[-(x - x_c)^2 / (2s_b^2)]}{\max_{z \in [0, 1]} \exp[-(z - x_c)^2 / (2s_b^2)]}, \quad x_c = 0.68, \quad s_b = 0.12. \quad (6)$$

Because $x_c \in [0, 1]$, the maximum is attained at $z = x_c$ and the denominator equals one analytically; it is retained to make the unit-peak convention explicit. In the finite-difference implementation, $b_j = b(x_j)$ at interior grid points, the two boundary entries are projected to $b_1 = b_n = 0$, and the interior vector is renormalized to $\max_j b_j = 1$. Thus the endpoints are a discrete boundary convention rather than zeros of the Gaussian itself.

The values $x_c = 0.68$ and $s_b = 0.12$ are fixed synthetic-benchmark design constants, chosen before any data are generated or models are trained. The center lies close to the $x = 0.72$ peak of the state-cost weight in Equation (9); the one-standard-deviation footprint has width $2s_b = 0.24$, preventing nearly pointwise actuation. The normalization leaves $a_t g_{\text{act}}$ responsible for the forcing amplitude. These constants are not statistical estimates and therefore have no associated estimation error in the present experiment. The present deployment shift instead varies the unobserved gain g_{act} . If actuator location or width were uncertain in an application, (x_c, s_b) would have to enter the latent parameter vector and the audit distribution; the guarantees reported here are conditional on the fixed actuator geometry.

For boundary lifting $\ell_b(x) = (1-x)b_L + xb_R$, define the affine state space $\mathcal{X}_B(b_L, b_R) = \ell_b + H_0^1(0, 1)$. Let $\Xi_B \subset (0, \infty) \times \mathbb{R}^2 \times \mathbb{R}$ denote the admissible parameter set for $\xi_B = (\nu, b_L, b_R, g_{\text{act}})$. The time- ΔT solution operator is the parameterized family

$$\mathcal{G}_B^{\Delta T}(\cdot, \cdot; \xi_B) : \mathcal{X}_B(b_L, b_R) \times \mathcal{A} \longrightarrow \mathcal{X}_B(b_L, b_R), \quad (u_t, a_t) \longmapsto u_{t+\Delta T}. \quad (7)$$

The world model observes $\xi_{B, \text{obs}} = (\nu, b_L, b_R)$ but not g_{act} . Varying this latent gain at deployment creates a controlled model shift.

Sampling and numerical solution. All Burgers transitions and trajectories are generated in this study; no external Burgers dataset is used. Each initial field is

$$u_0(x) = (1-x)b_L + xb_R + \sum_{k=1}^4 \frac{q_k z_k}{k} \sin(k\pi x), \quad q_k \sim \text{Unif}(q_{\min}, q_{\max}), \quad z_k \sim \mathcal{N}(0, 1). \quad (8)$$

Actions are sampled uniformly from $[-2, 2]$. Viscosity, both boundary values, initial-field amplitudes, and actuator gain are sampled independently from the intervals in Table 1. The plant uses 64 grid points, an upwind convective flux, centered diffusion, solver step 2×10^{-4} , and control interval 0.02.

The stage cost used for control is

$$\ell(u) + r(a) = \frac{1}{n} \sum_{j=1}^n w(x_j) u_j^2 + 0.002 a^2, \quad w(x) = 0.15 + 0.85 \exp \left[-\frac{1}{2} \left(\frac{x - 0.72}{0.14} \right)^2 \right]. \quad (9)$$

The source regime supplies proper training, validation, source calibration, and in-distribution testing. The deployment audit and principal shifted test use the combined-shift regime. Parameter-only and boundary-only tests isolate the two observed shift components.

Table 1: Sampling ranges for the synthetic Burgers data. Boundary entries are sampled independently from the listed interval.

Regime	ν	b_L, b_R	$(q_{\min}, q_{\max})$	g_{act}
Source	[0.015, 0.025]	[-0.03, 0.03]	[0.15, 0.55]	[0.60, 1.40]
Parameter shift	[0.006, 0.012]	[-0.03, 0.03]	[0.15, 0.55]	[0.60, 1.40]
Boundary shift	[0.015, 0.025]	[-0.16, 0.16]	[0.15, 0.55]	[0.60, 1.40]
Combined shift	[0.006, 0.012]	[-0.16, 0.16]	[0.30, 0.75]	[-0.50, 1.50]

3.2 Public two-dimensional Navier–Stokes data

The field-UQ stress test reuses the Reynolds-500 NeuralOperator benchmark [20, 31, 37]. We use `nsforcing_128.tgz` from the version-specific Zenodo record [doi:10.5281/zenodo.12825163](https://doi.org/10.5281/zenodo.12825163). Its spatial domain is the periodic torus $\mathbb{T}^2 = (0, 2\pi)^2$, and its time horizon is $T = 0.5$. In vorticity form, the dynamics are

$$\partial_t \omega + v \cdot \nabla \omega = \nu \Delta \omega - 4 \cos(4x_2), \quad \nabla \cdot v = 0, \quad \omega(x, 0) = \omega_0(x), \quad (x, t) \in \mathbb{T}^2 \times (0, 0.5]. \quad (10)$$

Here $\omega = \partial_{x_1} v_2 - \partial_{x_2} v_1$ and $v = \nabla^\perp \psi$, where the zero-mean stream function solves $-\Delta \psi = \omega$. The viscosity is fixed at $\nu = \text{Re}^{-1} = 2 \times 10^{-3}$. The archive therefore defines the flow map

$$\mathcal{G}_{\text{NS}}^T : L_0^2(\mathbb{T}^2) \longrightarrow L_0^2(\mathbb{T}^2), \quad \omega_0 \longmapsto \omega(\cdot, T), \quad (11)$$

where

$$L_0^2(\mathbb{T}^2) = \left\{ \omega \in L^2(\mathbb{T}^2) : \int_{\mathbb{T}^2} \omega(x) dx = 0 \right\} \quad (12)$$

is the space of mean-zero square-integrable periodic vorticity fields. The public input–output pairs contain no action channel. This operator is therefore used only to test high-dimensional field calibration.

3.3 Discretized benchmark dynamics and learned surrogates

We now connect each continuous solution operator to the finite array learned by its FNO surrogate. This notation makes explicit which claims concern function spaces and which concern one fixed grid [17, 21]. Let P_n denote evaluation on a fixed spatial grid. Let I_n denote a boundary-compatible reconstruction for Burgers and a periodic reconstruction for Navier–Stokes. For Burgers, $n = 64$ and the ground-truth one-step transition is

$$\begin{aligned} G_{B,n} : \mathbb{R}^{64} \times [-2, 2] \times \Xi_B &\longrightarrow \mathbb{R}^{64}, \\ G_{B,n}(x, a, \xi_B) &= P_n \mathcal{G}_B^{\Delta T}(I_n x, a; \xi_B). \end{aligned} \quad (13)$$

The residual FNO used by the controller approximates the partially observed map

$$\hat{G}_{\theta,B} : \mathbb{R}^{64} \times [-2, 2] \times \Xi_{B,\text{obs}} \longrightarrow \mathbb{R}^{64}. \quad (14)$$

Here $\Xi_{B,\text{obs}}$ contains the observed triples (ν, b_L, b_R) . Its training target is $y = G_{B,n}(x, a, (\xi_{B,\text{obs}}, g_{\text{act}}))$ even though the gain is not an input. The learned operator consequently represents the source-regime input–output relation rather than a gain-conditioned physical solver.

For Navier–Stokes, vectorizing the 128×128 fields gives $n = 16384$ and

$$\begin{aligned} G_{\text{NS},n} : \mathbb{R}^{16384} &\longrightarrow \mathbb{R}^{16384}, \\ G_{\text{NS},n}(x) &= P_n \mathcal{G}_{\text{NS}}^T(I_n x). \end{aligned} \quad (15)$$

The corresponding FNO has the same Euclidean input and output spaces, $\hat{G}_{\theta,\text{NS}} : \mathbb{R}^{16384} \rightarrow \mathbb{R}^{16384}$, and approximates $G_{\text{NS},n}$ on the distribution of archived vorticity fields. The term function-space set below means a joint set for this full discretized field, using the normalized geometry in Equation (1). It is not a continuum or resolution-transfer guarantee.

3.4 Data roles and split protocol

We distinguish four data roles:

$$\begin{aligned} \mathcal{D}_{\text{tr}} &= \{(x_i, a_i, y_i)\}_{i=1}^{N_{\text{tr}}}, & \mathcal{D}_{\text{val}} &= \{(x_i, a_i, y_i)\}_{i=1}^{N_{\text{val}}}, \\ \mathcal{D}_{\text{cal}} &= \{(x_i, a_i, y_i)\}_{i=1}^m, & \mathcal{D}_{\text{te}} &= \{(x_i, a_i, y_i)\}_{i=1}^{N_{\text{te}}}. \end{aligned}$$

Proper training and validation come from a source regime. The deployment audit $\mathcal{D}_{\text{cal}}$ is a small labeled sample from the target regime and is used only to compute conformal order statistics. The test set remains untouched until final evaluation. The clean and perturbed models, smoothing kernel, and floor rule are measurable with respect to proper-training randomness and are fixed before $\mathcal{D}_{\text{cal}}$ is scored.

This is an *audit-calibrated* setting: validity requires that calibration and target examples are exchangeable after conditioning on the trained models. If the deployment distribution changes again after the audit, ordinary split CP does not protect against that new shift. Weighted or adaptive CP [3], or covariate-shift calibration [35], is a possible extension, not a guarantee claimed here.

Table 2: Disjoint data roles in the final experiments. Burgers one-step test size is per regime. Trajectory calibration and test splits are separate from all one-step splits.

Benchmark	Train	Validation	Source cal.	Deploy. audit	Test	Traj. audit/test
Burgers	4000	600	800	300	800	200/400
Navier–Stokes	800	50	–	200	300	–

The release run uses random seed 27. Deterministic, role-specific seed offsets generate disjoint Burgers samples. For Navier–Stokes, proper-training, validation, audit, and test indices are non-overlapping. The clean FNO and the perturbed-label FNO receive the same proper-training inputs, while only the second model receives perturbed targets.

4 Methodology

4.1 Control-oriented dynamics ambiguity sets

The method maps a predictive scale to a state–action-dependent dynamics set $\mathcal{U}_\theta(x, a) \subset \mathbb{R}^n$ and exposes its geometry to a robust planner. We evaluate three properties separately. Predictive validity concerns the predeclared conformal containment event. Statistical efficiency concerns set width or volume. Control utility concerns paired mean cost, upper-tail cost, and control effort on independent closed-loop cases. A set can satisfy the first property without improving the third, so coverage is supporting evidence rather than a definition of control quality.

4.2 Residual FNO world model

The base one-step world model is a residual FNO,

$$\widehat{G}_\theta(x, a, \xi_{\text{obs}}) = x + \delta_\theta(x, a, \xi_{\text{obs}}). \quad (16)$$

For Burgers, the five input channels are the current field, the action multiplied by a fixed actuator profile, viscosity, a linear extension of the two boundary values, and the spatial coordinate. Each block applies a spectral convolution and a pointwise linear map inside a ResNet-style skip connection. The Fourier layer follows the FNO construction [20]; the identity skip follows residual learning [13]. Writing v_ℓ for a hidden field, one block has the form

$$v_{\ell+1} = v_\ell + \phi(W_\ell v_\ell + \mathcal{F}^{-1}(R_\ell \odot \mathcal{F} v_\ell)), \quad (17)$$

where R_ℓ retains the lowest Fourier modes and W_ℓ is a learned pointwise map. The FNO therefore combines global spectral mixing, a local linear path, and a residual connection. The final boundary coordinates are projected exactly to the supplied Dirichlet values. The full run uses width 32, 14 Fourier modes, four residual blocks, and 240,454 trainable parameters across the clean and perturbed operators.

FNO is chosen because it is a standard field propagator [20]; the broader framework is reviewed in [17]. Nothing in the calibration or robust-control construction requires Fourier layers. Replacing it by another differentiable operator changes the mean and scale quality, not the logic of the method.

4.3 Perturbation-based uncertainty scale

Following the stability principle used for data-scarce operator UQ [37], we train a second operator with the same architecture and the same proper-training inputs but perturbed labels:

$$\tilde{y}_i = y_i + \eta_i, \quad \eta_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, c^2 s_y^2 I_n). \quad (18)$$

Let $\tilde{G}_{\tilde{\theta}}$ be its prediction. The raw disagreement, smoothed disagreement, and stabilized scale are

$$d(x, a) = \left| \widehat{G}_\theta(x, a) - \tilde{G}_{\tilde{\theta}}(x, a) \right|, \quad (19)$$

$$\bar{d}(x, a) = K * d(x, a), \quad (20)$$

$$\sigma(x, a) = \bar{d}(x, a) \vee \tau_0. \quad (21)$$

The convolution is an odd moving-average window (five coordinates for Burgers; 15×15 for Navier–Stokes). On Burgers, τ_0 is 10% of the median interior disagreement on $\mathcal{D}_{\text{tr}}$, lower bounded by 10^{-6} . Smoothing prevents isolated near-zero denominators from dominating a max-type score; the floor is fixed before calibration. The raw scale has no coverage interpretation. Split conformal calibration supplies validity.

4.4 Perturbation-based nonconformity scores and calibration

Normalized residual scores are standard in adaptive conformal calibration [1], and the max-type operator score follows the perturbation-based construction of Wang et al. [37]. For residual $r = y - \widehat{G}_\theta(x, a)$, we compare three geometries.

$$S_{\text{L2}}(x, a, y) = \frac{\|r\|_{2,n}}{\|\sigma(x, a)\|_{2,n}}, \quad (22)$$

$$S_2(x, a, y) = \|r \odot \sigma(x, a)^{-1}\|_{2,n}, \quad (23)$$

$$S_\infty(x, a, y) = \|r \odot \sigma(x, a)^{-1}\|_\infty. \quad (24)$$

Here $\odot$ denotes coordinate-wise (Hadamard) multiplication and $\sigma(x, a)^{-1}$ denotes the vector of coordinate-wise reciprocals; the floor $\tau_0 > 0$ in Equation (21) makes every reciprocal finite. Applying Equation (3) separately gives q_{L2} , q_2 , and q_∞ . The corresponding one-step dynamics sets are

$$\mathcal{U}_{L2}(x, a) = \left\{ \hat{G}_\theta(x, a) + \Delta : \|\Delta\|_{2,n} \leq q_{L2} \|\sigma(x, a)\|_{2,n} \right\}, \quad (25)$$

$$\mathcal{U}_2(x, a) = \left\{ \hat{G}_\theta(x, a) + \Delta : \|\Delta \odot \sigma(x, a)^{-1}\|_{2,n} \leq q_2 \right\}, \quad (26)$$

$$\mathcal{U}_\infty(x, a) = \left\{ \hat{G}_\theta(x, a) + \Delta : |\Delta_j| \leq q_\infty \sigma_j(x, a), \ j = 1, \dots, n \right\}. \quad (27)$$

The first is an isotropic L^2 tube baseline. The second preserves heteroscedastic anisotropy but covers an L^2 event. The third targets simultaneous containment of every output coordinate. The box is usually wider because one failure anywhere in the field breaks the event.

4.5 Practical implementation

Algorithm 1 summarizes training, calibration, and deployment. The geometry index $g \in \{L2, 2, \infty\}$ selects one of Equations (22) to (24) and its associated set.

Algorithm 1 Perturbation-calibrated ambiguity sets for robust PDE control

Require: Disjoint $\mathcal{D}_{\text{tr}}, \mathcal{D}_{\text{val}}, \mathcal{D}_{\text{cal}}, \mathcal{D}_{\text{te}}$; miscoverage α ; noise scale c ; smoother K ; geometry g

Ensure: Frozen operators, multiplier q_g , and deployed set $\mathcal{U}_g$

- 1: $\theta \leftarrow \text{FITFNO}(\mathcal{D}_{\text{tr}} \text{ with clean targets}, \mathcal{D}_{\text{val}})$
 - 2: Form $\tilde{y}_i = y_i + \eta_i$ with $\eta_i \sim \mathcal{N}(0, c^2 s_y^2 I)$ on $\mathcal{D}_{\text{tr}}$
 - 3: $\tilde{\theta} \leftarrow \text{FITFNO}(\mathcal{D}_{\text{tr}} \text{ with } \tilde{y}_i, \mathcal{D}_{\text{val}})$
 - 4: Freeze $\hat{G}_\theta$ and $\tilde{G}_{\tilde{\theta}}$
 - 5: Fix K and τ_0 using proper-training predictions only
 - 6: **for** $(x_i, a_i, y_i) \in \mathcal{D}_{\text{cal}}$ **do**
 - 7: $\mu_i \leftarrow \hat{G}_\theta(x_i, a_i), \tilde{\mu}_i \leftarrow \tilde{G}_{\tilde{\theta}}(x_i, a_i)$
 - 8: $\sigma_i \leftarrow K * |\mu_i - \tilde{\mu}_i| \vee \tau_0$
 - 9: $r_i \leftarrow y_i - \mu_i$; compute $S_i^{(g)}$ from Equations (22) to (24)
 - 10: **end for**
 - 11: $k \leftarrow \lceil (|\mathcal{D}_{\text{cal}}| + 1)(1 - \alpha) \rceil$; $q_g \leftarrow S_{(k)}^{(g)}$
 - 12: Store weights, K , τ_0 , q_g , split metadata, and seed
 - Deployment:**
 - 13: **for** each MPC query (x_t, a) **do**
 - 14: Evaluate $\hat{G}_\theta(x_t, a)$ and $\sigma(x_t, a)$
 - 15: Construct $\mathcal{U}_g(x_t, a)$ and solve the robust inner problem
 - 16: **end for**
 - 17: Evaluate coverage and control once on $\mathcal{D}_{\text{te}}$
-

Checkpoint selection uses only $\mathcal{D}_{\text{val}}$. The smoothing rule, floor, and both operators are fixed before calibration, and no label in $\mathcal{D}_{\text{cal}}$ is reused for final coverage or closed-loop evaluation.

4.6 From one-step calibrated sets to multi-step robust control

Split conformal calibration constructs a set for one random transition, while MPC optimizes a counterfactual trajectory. The lifting from local prediction sets to trajectory uncertainty is therefore a separate modeling step. We use the robust-optimization principle of minimizing the worst-case objective over a declared uncertainty set [5]. For a candidate action sequence, define the admissible perturbations recursively as

$$\begin{aligned} \mathfrak{E}_H(x_0, a_{0:H-1}) &= \{\Delta_{0:H-1} : \Delta_t \in \mathcal{E}(x_t, a_t), \\ &\quad x_{t+1} = \hat{G}_\theta(x_t, a_t) + \Delta_t, \quad 0 \leq t < H\}, \end{aligned} \quad (28)$$

where $\mathcal{E}(x_t, a_t)$ is the centered ellipsoid or box from Equations (26) and (27). The robust planning problem is

$$\min_{a_{0:H-1} \in \mathcal{A}^H} \max_{\Delta_{0:H-1} \in \mathfrak{E}_H(x_0, a_{0:H-1})} \left\{ \sum_{t=1}^H \ell(x_t) + \sum_{t=0}^{H-1} r(a_t) \right\}. \quad (29)$$

Because $\mathcal{E}(x_t, a_t)$ changes with the perturbed state, $\mathfrak{E}_H$ is not generally a fixed Cartesian product. The implementation allows a separate disturbance choice at each predicted state but imposes no explicit temporal coupling beyond the state recursion. It is therefore a stagewise adversarial approximation. Its per-state disturbance choices are analogous to rectangular ambiguity in robust dynamic programming [15, 32], but the analogy is limited. Our sets depend on continuously valued perturbed FNO states, and we do not invoke the robust Bellman results of those papers. This computational choice is not an independence claim about physical model error. One-step marginal coverage does not provide joint coverage of Equation (28).

The remaining methodological problem is to construct trajectory-level sets that preserve temporal coherence and remain valid under policy-induced state and action distributions. Candidate extensions include a shared latent operator perturbation across the horizon, low-rank spatiotemporal sets, and trajectory conformal scores calibrated under a declared behavior or target policy. Their value should be assessed at matched trajectory coverage by set tightness and closed-loop cost, not by coverage alone.

4.6.1 Exact support and the adjoint robust counterpart

For the first-order robust counterpart, we freeze the local centered sets on the nominal rollout: $\mathcal{E}_t^0 = \mathcal{E}(\hat{x}_t, a_t)$ and $\mathcal{P}_H^0 = \prod_{t=0}^{H-1} \mathcal{E}_t^0$. This frozen approximation is a Cartesian product, so its support separates across time. This support-function view is standard in robust optimization [5].

Let λ_{t+1} denote the full derivative of the remaining nominal horizon cost with respect to x_{t+1} , expressed in the normalized inner product: $\delta J = \langle \lambda_{t+1}, \Delta_t \rangle_n$. Its value includes the effect of Δ_t on every later predicted state. The support of a centered set $\mathcal{E}$ is $h_{\mathcal{E}}(\lambda) = \sup_{\Delta \in \mathcal{E}} \langle \lambda, \Delta \rangle_n$.

Proposition 1 (Support functions of the calibrated geometries). *For positive scale σ ,*

$$h_{\mathcal{E}_2}(\lambda) = q_2 \|\lambda \odot \sigma\|_{2,n}, \quad (30)$$

$$h_{\mathcal{E}_{\infty}}(\lambda) = \frac{q_{\infty}}{n} \sum_{j=1}^n |\lambda_j| \sigma_j. \quad (31)$$

Proof. For the ellipsoid, set $\Delta = \sigma \odot z$ and apply Cauchy–Schwarz under $\langle \cdot, \cdot \rangle_n$; equality is attained by choosing z parallel to $\lambda \odot \sigma$. For the box, every coordinate is independent, and the maximizer is $\Delta_j = q_{\infty} \sigma_j \text{sign}(\lambda_j)$. $\square$

Linearizing the stacked inner problem over $\mathcal{P}_H^0$ gives

$$\hat{J}_H(a_{0:H-1}) + \sum_{t=0}^{H-1} h_{\mathcal{E}_t^0}(\lambda_{t+1}). \quad (32)$$

This formula is the decision interface: two sets can have similar coverage but produce different robust costs because their support differs along λ_{t+1} . Reverse-mode automatic differentiation supplies all finite-horizon adjoints in one backward pass [4].

4.6.2 Nonlinear adversarial rollout

The first-order objective can miss state-dependent scales and nonlinear error propagation. Our second inner solver parameterizes $\Delta_t = \sigma(x_t, a_t) \odot z_t$, initializes $z_{0:H-1} = 0$, performs projected gradient ascent on the full rollout cost, and projects each z_t onto either $\|z_t\|_{2,n} \leq q_2$ or $\|z_t\|_{\infty} \leq q_{\infty}$. For Burgers, $H = 8$ and $n = 64$, so the inner variable has $Hn = 512$ coordinates for each fixed action sequence. The objective is nonconvex because the rollout recursively composes nonlinear FNO blocks and recomputes $\sigma(x_t, a_t)$ at each perturbed state. Projected gradient search is a standard first-order primitive for constrained adversarial optimization [24]; its use here does not supply global optimality.

The experiment uses one zero initialization, three ascent steps, step size 0.8, and no random restarts. Projection makes the returned perturbation feasible for the stagewise sets. Its rollout cost is therefore a lower bound on the exact inner supremum, not an upper bound, global optimum, or certified optimality gap. We use this solver only as a compute-bounded adversarial-search baseline. The adjoint support and this nonlinear search answer different questions and neither certifies the full nonlinear min–max problem.

4.6.3 Outer optimization and receding-horizon deployment

The outer action sequence is optimized by the cross-entropy method (CEM): 128 candidates, 16 elites, four iterations, horizon eight, and action bound $|a_t| \leq 2$ [10]. The first action is applied to the numerical plant and the entire problem is replanned at the next observed state, as in receding-horizon MPC [28]. Nominal, isotropic-tube, adjoint-ellipsoid, adjoint-box, and adversarial controllers use the same world model and CEM budget.

4.7 Theoretical guarantees and scope

4.7.1 Finite-sample coverage on a fixed discretization

The next theorem is the standard split-conformal rank guarantee specialized to a discretized operator output [1, 19, 36]. We include the proof to make the conditioning and the simultaneous max-score event explicit.

Theorem 1 (One-step split-conformal coverage). *Condition on $\mathcal{D}_{\text{tr}}$, model-training randomness, perturbation noise, and all scale hyperparameters. If the m calibration triples and one new deployment triple are exchangeable, then each set constructed by its matching score satisfies*

$$\Pr \{G_{\star,n}(X, A) \in \mathcal{U}(X, A)\} \geq 1 - \alpha.$$

For $\mathcal{U}_\infty$, the event contains all n output coordinates simultaneously.

Proof. After conditioning, the score is fixed and identical for every example. The m calibration scores and the new score are therefore exchangeable. The rank argument for Equation (3), including the $+\infty$ convention, gives the result. For the max score, $S_\infty \leq q_\infty$ is equivalent to all n coordinate inequalities in Equation (27). $\square$

The theorem concerns one random transition. To obtain a different, valid trajectory statement, predeclare a behavior policy and define on whole length- H trajectories

$$S^{\text{traj}} = \max_{0 \leq t < H} \max_{1 \leq j \leq n} \frac{|x_{t+1,j} - \hat{x}_{t+1,j}|}{\sigma_j(\hat{x}_t, a_t)}. \quad (33)$$

Split-conformal calibration of this scalar score on exchangeable trajectories covers all times and coordinates of one new behavior-policy trajectory with probability at least $1 - \alpha$ by the same rank argument [19]. It does not certify counterfactual action sequences chosen by MPC.

4.7.2 First-order robustification and curvature

Let $\Delta = (\Delta_0, \dots, \Delta_{H-1})$ and let $F(\Delta)$ be the finite-horizon cost after injecting the stacked perturbation while the nominal-rollout sets remain frozen.

Proposition 2 (Conditional second-order remainder). *Suppose F is twice differentiable and $\|\nabla^2 F(\mathbf{z})\|_{\text{op}} \leq M_H$ on the line segments joining 0 to every $\Delta \in \mathcal{P}_H^0$. Then*

$$\sup_{\Delta \in \mathcal{P}_H^0} F(\Delta) \leq F(0) + \sum_{t=0}^{H-1} h_{\mathcal{E}_t^0}(\lambda_{t+1}) + \frac{M_H}{2} \sup_{\Delta \in \mathcal{P}_H^0} \|\Delta\|_2^2. \quad (34)$$

Proof. Taylor's theorem gives $F(\Delta) \leq F(0) + \langle \nabla F(0), \Delta \rangle + (M_H/2)\|\Delta\|_2^2$. The support of $\mathcal{P}_H^0$ is the sum of its component supports. Taking the supremum gives Equation (34). $\square$

The adjoint objective retains only the first two terms. A learned or audit-fitted quadratic feature is not a certified curvature bound unless a uniform M_H is established.

4.7.3 Multi-step propagation under uniform one-step error

Conformal coverage is distributional. The following result uses a different, deterministic assumption. Its recurrence is a deterministic specialization of the Lipschitz multi-step error mechanism analyzed for model-based reinforcement learning [2].

Assumption 1 (Uniform model error and Lipschitz dynamics). *On a forward-invariant state-action region $\mathcal{K}$,*

$$\left\| G_{\star,n}(x, a) - \hat{G}_\theta(x, a) \right\|_{2,n} \leq \epsilon, \quad \left\| G_{\star,n}(x, a) - G_{\star,n}(x', a) \right\|_{2,n} \leq L_G \|x - x'\|_{2,n}. \quad (35)$$

Proposition 3 (Open-loop rollout error). *For a common action sequence, equal initial states, and $e_t = \|x_t - \hat{x}_t\|_{2,n}$,*

$$e_{t+1} \leq L_G e_t + \epsilon, \quad e_h \leq \epsilon \sum_{k=0}^{h-1} L_G^k. \quad (36)$$

Proof. Add and subtract $G_{*,n}(\hat{x}_t, a_t)$, then use the two parts of Equation (35). Iterating the scalar recurrence with $e_0 = 0$ gives the geometric sum. $\square$

4.7.4 Discounted value and policy transfer

The value analysis follows the same Lipschitz propagation principle used to relate model error to value error in model-based reinforcement learning [2], specialized here to deterministic closed-loop maps and the normalized field norm.

For a feedback policy π , define the closed-loop map $F_G^\pi(x) = G(x, \pi(x))$ and stage reward $r_\pi(x) = r(x, \pi(x))$. The value is $V_G^\pi(x_0) = \sum_{t=0}^{\infty} \gamma^t r_\pi(x_t)$.

Theorem 2 (Fixed-policy value error). *Suppose the true and learned closed-loop maps for a fixed policy share the state Lipschitz constant L_G , their one-step discrepancy is uniformly at most ϵ , r_π is L_r -Lipschitz, and $\gamma L_G < 1$. Then*

$$\left| V_{G_*}^\pi(x_0) - V_{\hat{G}}^\pi(x_0) \right| \leq B(\epsilon) := \frac{\gamma L_r \epsilon}{(1 - \gamma)(1 - \gamma L_G)}. \quad (37)$$

Proof. By Equation (36), $e_t \leq \epsilon \sum_{k=0}^{t-1} L_G^k$. Hence

$$\sum_{t=1}^{\infty} \gamma^t L_r e_t \leq L_r \epsilon \sum_{k=0}^{\infty} L_G^k \sum_{t=k+1}^{\infty} \gamma^t = \frac{\gamma L_r \epsilon}{(1 - \gamma)(1 - \gamma L_G)}.$$

$\square$

Corollary 1 (Policy transfer with optimization error). *Suppose the preceding assumptions hold uniformly over a policy class. Let π^* maximize true-model value, and let $\hat{\pi}$ satisfy*

$$\sup_{\pi} V_{\hat{G}}^\pi(x_0) - V_{\hat{G}}^{\hat{\pi}}(x_0) \leq \delta_{\text{opt}}. \quad (38)$$

Then

$$V_{G_*}^{\pi^*}(x_0) - V_{\hat{G}}^{\hat{\pi}}(x_0) \leq \frac{2\gamma L_r \epsilon}{(1 - \gamma)(1 - \gamma L_G)} + \delta_{\text{opt}}. \quad (39)$$

If $L_G \leq 1$, the model term is at most $2\gamma L_r \epsilon / (1 - \gamma)^2$.

Proof. Insert and subtract the two learned-model values. The two model-transfer terms are bounded by $B(\epsilon)$, and learned-model suboptimality contributes δ_{opt} . $\square$

The CEM controller has finite sampling and four optimization iterations, so δ_{opt} is not known in our experiment. The corollary explains the target scaling; it is not a certificate for the reported CEM policy.

Numerical scaling witness. The deterministic rate can be checked without conformal calibration. For the scalar witness $G(x) = L_G x$, $\hat{G}(x) = L_G x + \epsilon$, and reward $-L_r |x|$, the fixed-policy bound is attained. Across six values each of ϵ and γ , the log-log slope is 1.000 in ϵ . After normalizing the value gap by $\gamma \epsilon$, the slope against effective horizon $(1 - \gamma)^{-1}$ is 2.000, with maximum relative numerical error 2.55×10^{-15} . Thus the calculation verifies the deterministic rate in Equation (39); it does not turn marginal conformal coverage into a uniform model-error guarantee.




Figure 2: Deterministic scaling audit for the policy-transfer result. Normalizing by $\gamma\epsilon$ exposes the predicted quadratic dependence on the effective discount horizon. The generic policy-transfer bound is twice the fixed-policy witness plus optimization error.

For ten controlled Burgers rollouts with an exactly injected one-step error, the maximum observed gap-to-finite-horizon-bound ratio is 0.359. On the learned FNO’s finite visited set, the maximum and mean gap-to-local-envelope ratios are 0.124 and 0.028. These are empirical local audits, not global PDE certificates.

4.7.5 Guarantee separation

Table 3: The three statement types used in the paper. None implies either of the others without additional assumptions.

Statement	Assumption	Event or quantity	What it does not provide
Split CP	Exchangeable audit/test scores	Marginal one-step field containment	Uniform error over a region or MPC safety
Trajectory CP	Exchangeable behavior-policy trajectories	All time/coordinates for one such trajectory	Counterfactual MPC action coverage
Value bound	Deterministic uniform ϵ , Lipschitz maps/reward	Rollout and policy-transfer error	A probability statement or conformal validity
Closed-loop test	Independent matched cases	Empirical mean/tail cost	A theorem beyond the tested design

5 Experimental evaluation

The numerical study follows the progression used in operator-learning conformal evaluations: model fit is checked first, predictive validity and set efficiency are then evaluated on held-out fields, and ablations vary one component while fixing the trained operators and test examples. Closed-loop control is reported only for the action-conditioned Burgers system. The Navier–Stokes data are used as a separate high-dimensional calibration test.

5.1 Baselines, metrics, and implementation

All Burgers controllers use the same residual FNO pair. We compare: uncontrolled dynamics; nominal MPC; source-calibrated isotropic L^2 robust MPC; deployment-audit L^2 robust MPC; audit ellipsoid with adjoint support; audit simultaneous box with adjoint support; and audit ellipsoid with nonlinear adversarial rollouts. This design separates access to deployment labels from ambiguity geometry and from the inner solver.

The predictive-set names encode two independent choices: where the conformal multiplier is calibrated and how the field residual is measured. All variants share the frozen FNO pair, the scale $\sigma(x, a)$, the target coverage

$1 - \alpha = 0.90$, and the same 800 compound-shift test transitions. *Source* L^2 uses the 800 source-regime calibration transitions with S_{L2} and the isotropic tube $\mathcal{U}_{L2}$ in Equations (22) and (25); it tests transfer under deployment shift. *Audit* L^2 uses the same score and set but re-estimates the multiplier on 300 target-regime audit transitions, thereby isolating the effect of calibration population. *Audit ellipsoid* uses that deployment audit with S_2 and $\mathcal{U}_2$ in Equations (23) and (26), combining coordinate-wise scales through one normalized L^2 score. *Audit simultaneous box* uses the same audit labels with S_∞ and $\mathcal{U}_\infty$ in Equations (24) and (27); its max score requires all grid coordinates to be covered in the same test field. Thus Source L^2 versus Audit L^2 isolates calibration shift, while the three audit rows isolate ambiguity geometry.

We report realized function-set coverage, decision-linear coverage (residual projected onto a cost direction), mean support, radius-error association, behavior-trajectory coverage, paired mean control cost, empirical $p90$ cost, and mean absolute action. Twenty-four combined-shift actuator-gain cases are matched across controllers, each with 20 receding-horizon steps. These cases share the same initial field and physical parameters except actuator gain; they are a controlled sweep rather than 24 independent PDE draws.

The two FNOs are trained for 60 epochs with AdamW, batch size 64, learning rate 2×10^{-3} , and weight decay 10^{-5} . The label-noise multiplier is $c = 0.05$. For every CEM candidate, the nonlinear adversary uses the single-start three-step projected-gradient budget specified above. Thus its reported row is an end-to-end controller result under a fixed search budget, not an evaluation of the exact inner maximum. The released package contains code, split rules, source JSON/CSV files, a portable notebook, and plotting scripts. Large raw data and full-run checkpoints are regenerated locally and are not stored in version control.

5.2 Ablation protocol and within-group controls

Following the controlled-comparison style of Wang et al. [37], every ablation group uses identical evaluation examples across variants and changes only the named interface. All reported groups condition on the seed-27 FNO pair; repeated cases quantify conditional variation, not training-seed variation. Table 4 records the comparison contract.

Table 4: Within-group ablation settings. “Fixed” lists the quantities shared by every row in the group; the final column is the only intended change.

Group	Fixed within the group	Varied factor
A: calibration and geometry	FNO pair, disagreement scale, 90% target, 800 ordered compound-shift transitions	Source versus deployment audit; L^2 tube versus ellipsoid versus coordinate box
B: robust-control interface	FNO pair, 24 matched plants, common initial field and physical parameters, 20 steps, common CEM seeds	Nominal, isotropic penalty, adjoint support, or single-start three-step PGD adversarial rollout
C: value-bound construction	Same 60 calibration and 160 test trajectories, observed gaps, horizon 20, $\gamma = 0.95$, 90% value target	Global recursion, local recursion, adjoint support, or adjoint plus curvature

Within Group A, source- L^2 versus audit- L^2 isolates calibration access; the three audit-calibrated rows isolate geometry. Within Group B, the two adjoint rows isolate geometry at a common first-order interface, whereas the ellipsoid-adjoint and adversarial-ellipsoid rows isolate the inner query at a fixed set. Within Group C, each raw bound is separately scaled on the common calibration trajectories before evaluation on the same independent test trajectories. We do not import the smoothing or label-noise ablation numbers from Wang et al.; those require retraining and are not claimed as our data.

5.3 World-model fit and uncertainty localization

The base FNO ends with training and validation MSEs 1.01×10^{-4} and 7.67×10^{-5} ; the perturbed replica has validation MSE 9.87×10^{-5} . Under combined shift, the mean field L^2 error is 0.00780, and the perturbation scale has a modest radius-error correlation of 0.275. The comparable validation losses show that label perturbation preserves the dominant one-step dynamics. The modest correlation, however, means that disagreement only weakly ranks where the base FNO will err. Conformal calibration can correct event frequency, but it cannot create coordinate-wise localization that is absent from the scale. An accurate mean predictor can therefore coexist with a wide calibrated set.

5.4 Ablation A: audit calibration restores coverage under shift

The first contrast in Table 5 changes only the calibration population. Source L^2 covers 593 of 800 compound-shift fields (0.741), 15.9 percentage points below the 0.90 target. Re-estimating the multiplier on the deployment audit,

without retraining either FNO, raises coverage to 707 of 800 (0.884). This paired change identifies calibration mismatch, rather than a change in model architecture, as the source of the largest coverage loss.

The second contrast changes only set geometry on the common deployment audit. The ellipsoid and simultaneous box each cover 725 of 800 fields (0.906), but they are not equivalent decision objects. The box attains 1.000 decision-linear coverage with mean support 0.04422, whereas the ellipsoid attains 0.986 with support 0.02012. The max-coordinate event protects every grid value and pays for directions that the current cost may not emphasize. Equal field coverage can therefore conceal more than a twofold difference in the uncertainty presented to the controller.

Table 5: Combined-shift one-step results. The target coverage is 0.90. Support is the half-width in the evaluated decision direction.

Set	Field coverage	Mean radius	Decision coverage	Mean support
Source L^2	0.741	0.01141	0.820	0.00650
Audit L^2	0.884	0.01882	0.934	0.01072
Audit ellipsoid	0.906	0.03242	0.986	0.02012
Audit simultaneous box	0.906	0.22729	1.000	0.04422




Figure 3: Realized one-step coverage across in-distribution, viscosity-shift, boundary-shift, and combined-shift test sets. Source calibration degrades most under combined shift; deployment-audit geometries recover the nominal region at different widths. Each regime contains 800 test transitions.

The separately calibrated behavior-trajectory max score attains empirical coverage 0.88 on 400 held-out length-10 trajectories at target 0.90. We report the observed proportion rather than rounding it to the target. Its scope is the fixed behavior-policy distribution in Equation (33). For scale, the audit- L^2 , ellipsoid, box, and trajectory events correspond to 707/800, 725/800, 725/800, and 352/400 covered examples. Their descriptive 95% Wilson intervals all contain 0.90. These intervals are an empirical stability check, not the source of the conformal guarantee.

5.5 Ablation B: matched-coverage sets yield different tail costs

The matched control sweep exposes the consequence of ambiguity geometry. The audit ellipsoid and simultaneous box have identical realized field coverage of 0.906, yet their adjoint controllers reach empirical p_{90} costs of 1.0201 and 0.9586. Relative to nominal MPC, the box lowers mean cost by 1.89% and p_{90} by 10.84%; audit L^2 lowers them by 1.92% and 7.79%. All robust variants improve the empirical tail, while every mean improvement remains below 2%. Robustification therefore changes the difficult actuator-gain cases more than the center of the matched sweep.

The box-adjoint controller also reduces mean $|a|$ from 1.278 to 1.101. This association is consistent with less aggressive actions under latent gain uncertainty, but the experiment does not identify action moderation as the sole cause of the cost reduction. The nonlinear adversarial ellipsoid is not best under the single-start three-step PGD and fixed CEM budgets. This outcome does not establish that nonlinear min–max robustification is inferior. It may instead reflect local optimization error in the high-dimensional nonconvex inner search. The comparison therefore concerns the implemented query mechanisms under stated budgets, not their globally optimized counterparts. Because the 24 cases form one matched gain sweep, these percentages are conditional effect sizes rather than population-level estimates across training seeds.

Table 6: Closed-loop cost over 24 matched actuator-gain cases. Changes are relative to nominal MPC; lower is better.

Controller	Mean	Mean change	$p90$	$p90$ change	Mean $ a $
Nominal MPC	0.5797	–	1.0752	–	1.278
Source- L^2 robust	0.5724	–1.27%	1.0250	–4.66%	1.185
Audit- L^2 robust	0.5686	–1.92%	0.9914	–7.79%	1.115
Ellipsoid-adjoint	0.5747	–0.88%	1.0201	–5.12%	1.230
Box-adjoint	0.5688	–1.89%	0.9586	–10.84%	1.101
Adversarial ellipsoid	0.5764	–0.58%	1.0532	–2.04%	1.262




Figure 4: Empirical $p90$ closed-loop cost for the matched combined-shift actuator-gain sweep. The simultaneous-box adjoint controller gives the largest tail reduction, while the nonlinear adversarial controller is not best in this finite compute budget. Its row uses one zero start and three projected gradient-ascent steps, so it is not a globally optimized min–max baseline.

5.6 Navier–Stokes stress test: validity survives, localization weakens

On the official 128×128 test set, the base FNO has mean field RMSE 0.1635, where each sample uses the square root of its mean coordinate-wise squared error. This quantity is not relative L^2 error. The max-type multiplier is $q_{0.90} = 30.679$; simultaneous test coverage is 0.883 and mean full coordinate-band width is 4.3409. The pointwise disagreement scale correlates only $r = 0.169$ with absolute error. The realized count is 265/300, with descriptive 95% Wilson interval $[0.842, 0.915]$, which contains the 0.90 target. Thus the observed 0.883 does not by itself diagnose systematic undercoverage at this sample size.

The multiplier and width reveal the more important limitation. A max score over 16,384 coordinates amplifies the largest standardized residual, while $r = 0.169$ shows that the local scale ranks pointwise errors weakly. Calibration

then obtains approximate simultaneous validity mainly through a large global multiplier rather than accurate spatial adaptation. This negative result distinguishes a valid conformal wrapper from an informative uncertainty model: calibration can correct the event frequency, but it cannot create localization that is absent from the underlying scale.

Table 7: Independent-test Navier–Stokes field-UQ results. No action channel is available, so no control metric is reported.

Field RMSE	$q_{0.90}$	Simultaneous coverage	Full width	Scale-error r
0.1635	30.679	0.883	4.3409	0.169




Figure 5: Representative absolute error of the two-dimensional FNO prediction. Errors are spatially structured rather than uniform, motivating a local scale; the following quantitative association test shows that the learned scale only weakly tracks them.




Figure 6: Pointwise perturbation scale versus absolute error on the independent Navier–Stokes test split. Pearson $r = 0.169$ reveals weak localization and helps explain the large simultaneous conformal multiplier.

5.7 Ablation C: tighter value bounds require matched coverage

Four raw bound constructions are scaled on the same 60 calibration trajectories to the 90% value-level target and evaluated on 160 independent trajectories at horizon 20 and $\gamma = 0.95$. The local recursion reaches 0.944 coverage with mean bound 0.0870. It is 11.9% narrower than the global recursion while attaining slightly higher empirical coverage. This is the only comparison that supports a clear reduction in conservatism at matched validity.

Adjoint support is wider, with mean bound 0.2035, and has lower coverage of 0.925. The fitted curvature coefficient is zero, so adding the curvature feature leaves the last two rows identical. Utilization above one corresponds to an uncovered test trajectory and cannot be read as tightness. The result shows that an adjoint direction can be useful for choosing robust actions without yielding the tightest scalar certificate of the realized value gap.

Table 8: Independent-test bound comparison. Utilization is observed value gap divided by the reported bound. Coverage must be read together with width.

Method	Coverage	Mean bound	Median η	$p90 \ \eta$	Max η
Global maximum	0.938	0.0987	0.210	0.688	3.919
Local recursion	0.944	0.0870	0.322	0.788	2.190
Adjoint support	0.925	0.2035	0.265	0.854	1.766
Adjoint + curvature	0.925	0.2035	0.265	0.854	1.766




Figure 7: Coverage–mean-bound frontier on independent trajectories. A bound is less conservative only when it is narrower at matched coverage; utilization alone can be inflated by undercoverage.

5.8 Interpretation and limitations

What the current evidence supports. Three separations organize the evidence. Calibration population changes validity under compound shift; geometry changes support along the MPC adjoint; and the inner query changes how that support enters an action. In the controlled actuator sweep, the distinction is largest in the upper tail. The simultaneous box is wide in field space yet gives the lowest empirical $p90$ cost. Predictive width and decision utility are therefore not the same ordering.

What the current evidence does not support. The Burgers control result is one seed and one matched actuator sweep, not a multi-seed population claim. The public Navier–Stokes pairs have no action channel. The ordinary split-conformal theorem requires exchangeability with the deployment audit and does not cover a later distribution change. Marginal one-step coverage does not certify the stagewise uncertainty model queried by MPC. The behavior-trajectory result also does not cover counterfactual optimized actions. The deterministic value theorem assumes a uniform error and common closed-loop Lipschitz constants that were not certified for the learned FNO. The adjoint robust counterpart is first order unless a uniform curvature constant is established. At $H = 8$ and $n = 64$, the nonlinear inner problem has 512 perturbation coordinates for each candidate action sequence. Its state-dependent FNO rollout is nonconvex. Single-start three-step PGD supplies only a feasible lower bound on the worst-case cost, with no convergence test, multi-start check, or dual upper bound. The returned action therefore has no certificate against all perturbations in the stagewise set. Finite CEM adds a second optimization error that is also measured empirically rather than bounded.

Negative results guide the next experiment. The nonlinear adversarial controller is not best; the adjoint scalar bound is not tighter than the local recursion; the curvature feature adds nothing; and the Navier–Stokes perturbation scale has weak error association. These outcomes argue for the following next tests: multiple training/control seeds; audit-size and shift-severity sweeps; learned low-rank covariance or spectral function-space sets; policy-coupled trajectory calibration; and a controlled two-dimensional PDE with an explicit action channel. The adversarial audit also requires iteration and step-size sweeps, random restarts, and reporting the best feasible objective against the adjoint approximation. Deep ensembles, probabilistic neural operators, and learned quantile operators should be added as scale baselines under a matched total label and compute budget [18], probabilistic neural operators [7], and calibrated quantile operators [23].

6 Conclusion

Predictive coverage is not yet a control interface. This work makes that interface explicit by calibrating a field-valued dynamics set and querying its geometry along directions that change a finite-horizon objective. Exact support functions and adversarial rollouts provide two realizations, while a guarantee ledger keeps one-step conformal validity, deterministic value propagation, and empirical closed-loop performance separate.

The decisive controlled-Burgers comparison observes identical realized field coverage of 0.906 while changing geometry: the ellipsoid and box lead to empirical $p90$ costs of 1.0201 and 0.9586. This is conditional evidence from one matched actuator sweep, not a population or safety guarantee. The Navier–Stokes stress test exposes the complementary limitation, since high-dimensional validity can require wide sets when the local scale is weak. The next method should be held to a harder standard: policy-level trajectory coverage together with lower closed-loop cost at matched coverage and compute.

Data availability

The controlled-Burgers data are synthetic data generated in this study by the released finite-difference simulator using the equations, parameter ranges, split sizes, and random seeds specified in the manuscript and repository; no external Burgers dataset was used. The derived metrics and figure-source tables are included in the project repository. The Navier–Stokes experiments reuse the public NeuralOperator Team Reynolds-500 dataset, specifically `nsforcing_128.tgz`, available from Zenodo at [doi:10.5281/zenodo.12825163](https://doi.org/10.5281/zenodo.12825163) [31]. The external archive is not redistributed in this repository and should be obtained from the cited record.

Code availability

The simulator, FNO training and calibration code, robust-control experiments, notebook, derived result tables, and plotting scripts are available at <https://github.com/YaowenKANG666/decision-calibrated-pde-control>.

References

- [1] Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023.
- [2] Kavosh Asadi, Dipendra Misra, and Michael Littman. Lipschitz continuity in model-based reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 264–273, 2018.
- [3] Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Conformal prediction beyond exchangeability. *The Annals of Statistics*, 51(2):816–845, 2023.
- [4] Atilim Gunes Baydin, Barak A. Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. Automatic differentiation in machine learning: A survey. *Journal of Machine Learning Research*, 18(153):1–43, 2018.
- [5] Dimitris Bertsimas, David B. Brown, and Constantine Caramanis. Theory and applications of robust optimization. *SIAM Review*, 53(3):464–501, 2011.
- [6] Nicolas Boullé and Alex Townsend. A mathematical guide to operator learning. In *Handbook of Numerical Analysis*, volume 25, pages 83–125. Elsevier, 2024.
- [7] Christopher Bülte, Philipp Scholl, and Gitta Kutyniok. Probabilistic neural operators for functional uncertainty quantification. *Transactions on Machine Learning Research*, 2025. arXiv:2502.12902.
- [8] Kong Yao Chee, Thales C. Silva, M. Ani Hsieh, and George J. Pappas. Uncertainty quantification and robustification of model-based controllers using conformal prediction. In *Proceedings of the 6th Annual Learning for Dynamics and Control Conference*, volume 242 of *Proceedings of Machine Learning Research*, pages 528–540, 2024.
- [9] Jose Leopoldo Contreras, Ola Shorinwa, and Mac Schwager. Safe, out-of-distribution-adaptive MPC with conformalized neural network ensembles. In *Proceedings of the 7th Annual Learning for Dynamics and Control Conference*, volume 283 of *Proceedings of Machine Learning Research*, pages 194–207, 2025.
- [10] Pieter-Tjerk de Boer, Dirk P. Kroese, Shie Mannor, and Reuven Y. Rubinstein. A tutorial on the cross-entropy method. *Annals of Operations Research*, 134(1):19–67, 2005.
- [11] Anushri Dixit, Lars Lindemann, Skylar X. Wei, Matthew Cleaveland, George J. Pappas, and Joel W. Burdick. Adaptive conformal prediction for motion planning among dynamic agents. In *Proceedings of the 5th Annual Learning for Dynamics and Control Conference*, volume 211 of *Proceedings of Machine Learning Research*, pages 300–314, 2023.
- [12] Amir-massoud Farahmand, Andre Barreto, and Daniel Nikovski. Value-aware loss function for model-based reinforcement learning. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54 of *Proceedings of Machine Learning Research*, pages 1486–1494, 2017.
- [13] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.
- [14] Peiyan Hu, Xiaowei Qian, Wenhao Deng, Rui Wang, Haodong Feng, Ruiqi Feng, Tao Zhang, Long Wei, Yue Wang, Zhi-Ming Ma, and Tailin Wu. From uncertain to safe: Conformal adaptation of diffusion models for safe PDE control. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 24653–24677, 2025.
- [15] Garud N. Iyengar. Robust dynamic programming. *Mathematics of Operations Research*, 30(2):257–280, 2005.
- [16] Chancellor Johnstone and Bruce Cox. Conformal uncertainty sets for robust optimization. In *Proceedings of the Tenth Symposium on Conformal and Probabilistic Prediction and Applications*, volume 152 of *Proceedings of Machine Learning Research*, pages 72–90, 2021.
- [17] Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces with applications to PDEs. *Journal of Machine Learning Research*, 24(89):1–97, 2023.

- [18] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, volume 30, pages 6402–6413, 2017.
- [19] Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- [20] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations*, 2021.
- [21] Lars Lindemann, David Millard, and Ali Baheri. Split conformal prediction in the function space with neural operators. In *International Conference on Learning Representations*, 2026.
- [22] Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3(3):218–229, 2021.
- [23] Ziqi Ma, David Pitt, Kamyar Azizzadenesheli, and Anima Anandkumar. Calibrated uncertainty quantification for operator learning via conformal prediction. *Transactions on Machine Learning Research*, 2024.
- [24] Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards deep learning models resistant to adversarial attacks. In *International Conference on Learning Representations*, 2018.
- [25] Emilia Magnani, Nicholas Krämer, Runa Eschenhagen, Lorenzo Rosasco, and Philipp Hennig. Approximate Bayesian neural operators: Uncertainty quantification for parametric PDEs. *Transactions on Machine Learning Research*, 2025. arXiv:2208.01565.
- [26] Emilia Magnani, Marvin Pförtner, Tobias Weber, and Philipp Hennig. Linearization turns neural operators into function-valued Gaussian processes. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- [27] Ali Malik, Volodymyr Kuleshov, Jiaming Song, Danny Nemer, Harlan Seymour, and Stefano Ermon. Calibrated model-based deep reinforcement learning. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 4314–4323, 2019.
- [28] David Q. Mayne, James B. Rawlings, Christopher V. Rao, and Pierre O. M. Scokaert. Constrained model predictive control: Stability and optimality. *Automatica*, 36(6):789–814, 2000.
- [29] S. Chandra Mouli, Danielle C. Maddix, Shima Alizadeh, Gaurav Gupta, Andrew Stuart, Michael W. Mahoney, and Bernie Wang. Using uncertainty quantification to characterize and improve out-of-domain learning for PDEs. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 36372–36418, 2024.
- [30] Christian Moya, Ali Mollaali, Zhen Zhang, Lu Lu, and Guang Lin. Conformalized-DeepONet: A distribution-free framework for uncertainty quantification in deep operator networks. *Physica D: Nonlinear Phenomena*, 471:134418, 2025.
- [31] NeuralOperator Team. Navier–Stokes dataset. Zenodo, 2024. Version-specific record DOI: 10.5281/zenodo.12825163.
- [32] Arnab Nilim and Laurent El Ghaoui. Robust control of Markov decision processes with uncertain transition matrices. *Operations Research*, 53(5):780–798, 2005.
- [33] Anutam Srinivasan, Antoine Leeman, and Glen Chou. Safety beyond the training data: Robust out-of-distribution MPC via conformalized system level synthesis. In *Proceedings of the 8th Annual Learning for Dynamics and Control Conference*, volume 331 of *Proceedings of Machine Learning Research*, pages 412–439, 2026.
- [34] Shyam Sundhar Ramesh, Pier Giuseppe Sessa, Yifan Hu, Andreas Krause, and Ilija Bogunovic. Distributionally robust model-based reinforcement learning with large state spaces. In *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pages 100–108, 2024.
- [35] Ryan J. Tibshirani, Rina Foygel Barber, Emmanuel J. Candès, and Aaditya Ramdas. Conformal prediction under covariate shift. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- [36] Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, New York, 2005.
- [37] Weinan Wang, Bowen Gang, and Hao Deng. Operator learning for the 2D incompressible Navier–Stokes equations: A conformal prediction approach in the data-scarce regime, 2026.